



Replica sets

<http://www.mongodb.org/display/DOCS/Replica+Sets>

Replica sets are basically master/slave replication with automatic failover. The current master is called the **primary** and the servers syncing off of it are called **secondaries**. If a primary becomes unreachable, a secondary may be elected to be primary without human intervention.

A majority of servers must be able to reach the primary, both to elect it and to keep it a primary.

What is a majority?

If your set consists of:

- 1 server, 1 server is a majority
- 2 servers, 2 servers are a majority
- 3 servers, 2 servers are a majority
- 4 servers, 3 servers are a majority
- ...

If you have two servers and you want one server to be elected primary if the other goes down, you can use an *arbiter*, a lightweight *mongod* process that only participates in elections. You need at most 1 arbiter.

Config Format

```
{
  "_id" : "nameOfSet"
  "members" : [
    { "_id" : 0, "host" : "host0" },
    { "_id" : 1, "host" : "host1" },
    ...
    { "_id" : n, "host" : "hostn" }
  ]
}
```

Member objects can also contain:

- `priority : n` ➤ Members will be elected primary in order of priority, if possible.
n=0 means this member will never be master
- `slaveDelay : n` ➤ This member will always be a secondary and will lag *n* seconds behind the master
- `arbiterOnly : true` ➤ This member will be an arbiter
- `hidden : true` ➤ Do not show this member in `isMaster` output
- `tags : [...]` ➤ Member location description, see <http://www.mongodb.org/display/DOCS/Data+Center+Awareness>

Shell Helpers

`rs.initiate()` ➤ Creates a new replica set with one member

`rs.add("host:port")` ➤ Adds a member

`rs.add({"_id" : n, "host" : "host:port", slaveDelay : seconds})` ➤ Adds a member with non-default properties

`rs.addArb("host:port")` ➤ Adds an arbiter

`rs.remove("host:port")` ➤ Removes a member

Administration

Please review the full documentation for more details about administration of replica sets.

`rs.status()` ➤ See the status of each member

`rs.conf()` ➤ See the current set configuration

`rs.reconfig(newConfig)` ➤ Change the set configuration

`rs.isMaster()` ➤ See which member is primary

`rs.stepDown(n)` ➤ Force the primary to become a secondary for *n* seconds

`rs.freeze(n)` ➤ Prevent a secondary from becoming the primary for *n* seconds. (*n* = 0 means unfreeze.)

Sharding

To add a replica set as a shard, run:

```
db.adminCommand({addShard : "setName/host:port[,host2:port2]"})
```

mongos will use the given host(s) as seeds and find the rest of the set. It will also automatically detect when a replica set's members change.

Tips

- Run one secondary with journaling enabled to ensure that you always have a clean copy of your data available.
- Start replica set members with the `--rest` option for useful information in the web interface.